

BLG 454E Learning from Data

FALL 2022-2023



Assoc. Prof. Yusuf Yaslan

Trees

Lecture Notes from Alpaydm 2010 Introduction to Machine Learning 2e © The MIT Press AND
S. Theodoridis, Machine Learning: A Bayesian and Optimization Perspective, Elsevier 2nd Edition AND
P. Flach, Chapter 5 Trees, Machine Learning: Making Sense of Data 2013 AND
N. De Freitas, Machine Learning-Decision Trees, Random Forest, 2013 Youtube Lecture AND
A. Criminisi et al. Decision Forests for Classification, Regression, Density Estimation, Manifold Learning and
Semi-Supervised Learning, 2011, Microsoft Technical Report

Trees/Decision Trees

- Each internal node is labelled with a feature
- Each edge is labelled with a literal
 - Set of literals: Split
- Leaves are
 - Classification: class labels
 - Regression: numeric, average, or local fit



Trees/Decision Trees

- Univariate case
 - Single attribute x_1
 - In case of numeric, binary split $x_1 > m$
 - In case of discrete, n-way split
- Multivariate case
 - All attributes, x
- Learning is *greedy*
 - *Finding the best split (and feature) recursively*

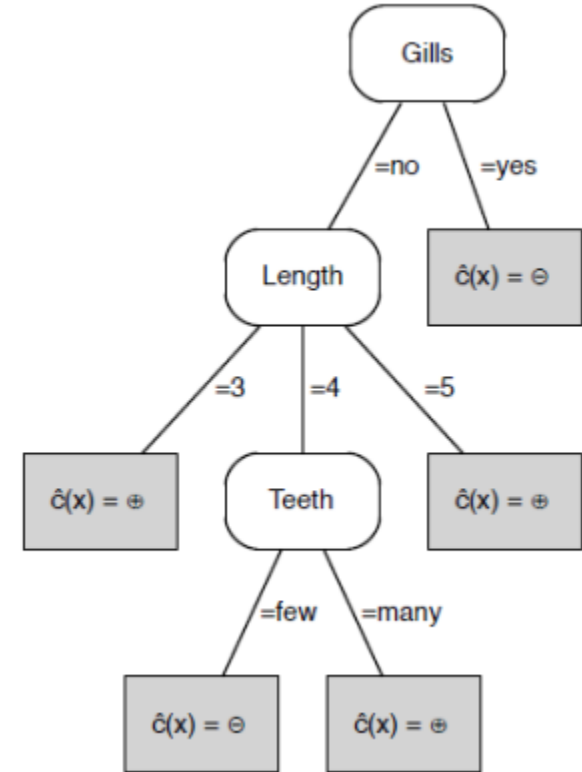


Figure from P. Flach, Chapter 5 Trees, Machine Learning: Making Sense of Data 2013

Remarks on Trees

- + Trees can naturally treat mixtures of numeric and categorical variables.
- + They scale well with large data sets.
- + They can treat missing variables in an effective way.
- + They are easily interpretable (expressive).
- Prediction performance may not be as good as other classifiers such as SVM or NNs
- Trees are sensitive to changes in training data – unstable
- **Bagging/Boosting** can reduce variance and improve generalization performance
- **Random forests** use bagging and often have very good prediction accuracy.


Classification Trees (ID3, CART, C4.5)

- For node m , N_m instances reach m , N_m^i belong to C_i

$$\hat{P}(C_i | \mathbf{x}, m) \equiv p_m^i = \frac{N_m^i}{N_m}$$

- Node m is **pure** if p_m^i is 0 or 1
- Measure of **impurity** is **entropy**

$$I_m = -\sum_{i=1}^K p_m^i \log_2 p_m^i$$



Lecture Notes for E Alpaydm 2004 Introduction to Machine Learning © The MIT Press (V1.1)

Impurity

- Every split from X_t to X_{tL} and X_{tR} must generate class-homogeneous sets compared to X_t .
 - Easier to distinguish classes in the new subsets.
- Choose the feature that maximizes the decrease in node impurity
 - Higher **Information Gain**

$$I = H(\mathcal{S}) - \sum_{i \in \{1,2\}} \frac{|\mathcal{S}^i|}{|\mathcal{S}|} H(\mathcal{S}^i)$$

- Besides *entropy*
 - *Gini Index*:

$$I_m = - \sum_{i=1}^K p_m^i (1 - p_m^i)$$

Suppose we have the following five positive examples (the first three are the same as in [Example 4.1](#)):

p1: Length = 3 \wedge Gills = no \wedge Beak = yes \wedge Teeth = many

p2: Length = 4 \wedge Gills = no \wedge Beak = yes \wedge Teeth = many

p3: Length = 3 \wedge Gills = no \wedge Beak = yes \wedge Teeth = few

p4: Length = 5 \wedge Gills = no \wedge Beak = yes \wedge Teeth = many

p5: Length = 5 \wedge Gills = no \wedge Beak = yes \wedge Teeth = few

and the following negatives (the first one is the same as in [Example 4.2](#)):

n1: Length = 5 \wedge Gills = yes \wedge Beak = yes \wedge Teeth = many



n2: Length = 4 \wedge Gills = yes \wedge Beak = yes \wedge Teeth = many

n3: Length = 5 \wedge Gills = yes \wedge Beak = no \wedge Teeth = many

n4: Length = 4 \wedge Gills = yes \wedge Beak = no \wedge Teeth = many

n5: Length = 4 \wedge Gills = no \wedge Beak = yes \wedge Teeth = few

Example



Similar calculations for Beak and Teeth would give 0.76 and 0.97 respectively.

Pick 'Gills' as the first feature!

Length = [3, 4, 5]	[2+, 0-][1+, 3-][2+, 2-]
Gills = [yes, no]	[0+, 4-][5+, 1-]
Beak = [yes, no]	[5+, 3-][0+, 2-]
Teeth = [many, few]	[3+, 4-][2+, 1-]

Impurity of the root is 5+, 5- so it is 1.

We will find the feature that maximizes the decrease in impurity.

Impurity of using 'Length' is

$$-2/2 * \log 1 - 0/2 * \log 0/2 = 0$$

$$-1/4 * \log(1/4) - 3/4 * \log(3/4) = 0.81$$

$$-2/4 * \log(2/4) - 2/4 * \log(2/4) = 1$$

Total entropy is weighted average:

$$2/10 * 0 + 4/10 * 0.81 + 4/10 * 1 = 0.72$$

Impurity of 'Gills' is

$$-0/4 * \log 0 - 4/4 * \log 1 = 0$$

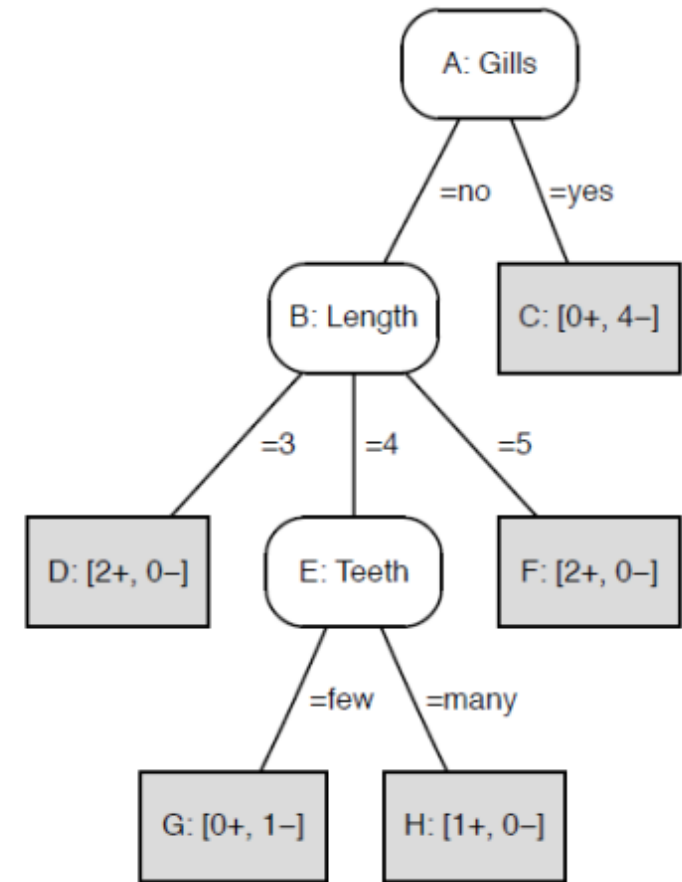
$$-5/6 * \log(5/6) - 1/6 * \log(1/6) =$$

Total entropy is

$$4/10 * 0 + 6/10 * (-5/6 * \log(5/6) - 1/6 * \log(1/6)) = 0.39$$

Example cont'd

- After the first feature is selected, the right side of the tree reaches the final leaf node. NEGATIVE
- Perform similar operations to decide on the second feature, then third, and forth until you reach the class labels in the leaf nodes.



Another Example

- Let's compare *Pat* (Patrons) and *Type*
 - $I(\text{Pat}) = 0.541$, $I(\text{Type}) = 0$

Example	Input Attributes										Goal
	<i>Alt</i>	<i>Bar</i>	<i>Fri</i>	<i>Hun</i>	<i>Pat</i>	<i>Price</i>	<i>Rain</i>	<i>Res</i>	<i>Type</i>	<i>Est</i>	<i>WillWait</i>
x ₁	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>Yes</i>	<i>Some</i>	<i>\$\$\$</i>	<i>No</i>	<i>Yes</i>	<i>French</i>	<i>0–10</i>	$y_1 = \text{Yes}$
x ₂	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>Yes</i>	<i>Full</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Thai</i>	<i>30–60</i>	$y_2 = \text{No}$
x ₃	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>Some</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Burger</i>	<i>0–10</i>	$y_3 = \text{Yes}$
x ₄	<i>Yes</i>	<i>No</i>	<i>Yes</i>	<i>Yes</i>	<i>Full</i>	<i>\$</i>	<i>Yes</i>	<i>No</i>	<i>Thai</i>	<i>10–30</i>	$y_4 = \text{Yes}$
x ₅	<i>Yes</i>	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>Full</i>	<i>\$\$\$</i>	<i>No</i>	<i>Yes</i>	<i>French</i>	<i>>60</i>	$y_5 = \text{No}$
x ₆	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>Yes</i>	<i>Some</i>	<i>\$\$</i>	<i>Yes</i>	<i>Yes</i>	<i>Italian</i>	<i>0–10</i>	$y_6 = \text{Yes}$
x ₇	<i>No</i>	<i>Yes</i>	<i>No</i>	<i>No</i>	<i>None</i>	<i>\$</i>	<i>Yes</i>	<i>No</i>	<i>Burger</i>	<i>0–10</i>	$y_7 = \text{No}$
x ₈	<i>No</i>	<i>No</i>	<i>No</i>	<i>Yes</i>	<i>Some</i>	<i>\$\$</i>	<i>Yes</i>	<i>Yes</i>	<i>Thai</i>	<i>0–10</i>	$y_8 = \text{Yes}$
x ₉	<i>No</i>	<i>Yes</i>	<i>Yes</i>	<i>No</i>	<i>Full</i>	<i>\$</i>	<i>Yes</i>	<i>No</i>	<i>Burger</i>	<i>>60</i>	$y_9 = \text{No}$
x ₁₀	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Full</i>	<i>\$\$\$</i>	<i>No</i>	<i>Yes</i>	<i>Italian</i>	<i>10–30</i>	$y_{10} = \text{No}$
x ₁₁	<i>No</i>	<i>No</i>	<i>No</i>	<i>No</i>	<i>None</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Thai</i>	<i>0–10</i>	$y_{11} = \text{No}$
x ₁₂	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Yes</i>	<i>Full</i>	<i>\$</i>	<i>No</i>	<i>No</i>	<i>Burger</i>	<i>30–60</i>	$y_{12} = \text{Yes}$

Figure 18.3 Examples for the restaurant domain.

Regression Trees

- Error at node m :

$$b_m(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{x} \in \mathcal{X}_m : \mathbf{x} \text{ reaches node } m \\ 0 & \text{otherwise} \end{cases}$$





$$E_m = \frac{1}{N_m} \sum_t (r^t - g_m)^2 b_m(\mathbf{x}^t) \quad g_m = \frac{\sum_t b_m(\mathbf{x}^t) r^t}{\sum_t b_m(\mathbf{x}^t)}$$

- After splitting:

$$b_{mj}(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{x} \in \mathcal{X}_{mj} : \mathbf{x} \text{ reaches node } m \text{ and branch } j \\ 0 & \text{otherwise} \end{cases}$$

$$E'_m = \frac{1}{N_m} \sum_j \sum_t (r^t - g_{mj})^2 b_{mj}(\mathbf{x}^t) \quad g_{mj} = \frac{\sum_t b_{mj}(\mathbf{x}^t) r^t}{\sum_t b_{mj}(\mathbf{x}^t)}$$

Model Selection in Trees:

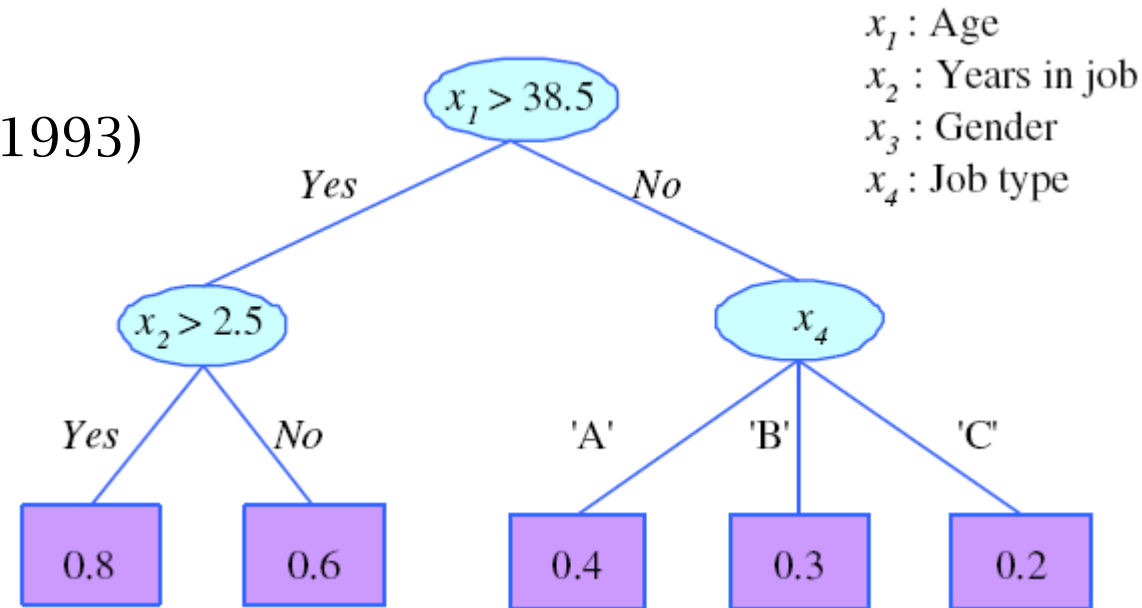


Pruning Trees

- Remove subtrees for better generalization (decrease variance)
 - Prepruning: Early stopping
 - Postpruning: Grow the whole tree then prune subtrees which overfit on the pruning set
- Prepruning is faster, postpruning is more accurate (requires a separate pruning set)

Rule Extraction from Trees

C4.5Rules
(Quinlan, 1993)



- R1: IF (age>38.5) AND (years-in-job>2.5) THEN $y = 0.8$
R2: IF (age>38.5) AND (years-in-job \leq 2.5) THEN $y = 0.6$
R3: IF (age \leq 38.5) AND (job-type='A') THEN $y = 0.4$
R4: IF (age \leq 38.5) AND (job-type='B') THEN $y = 0.3$
R5: IF (age \leq 38.5) AND (job-type='C') THEN $y = 0.2$

Random Trees

Random Forest

- How to build a random tree?
 - Randomly select two features, pick the best one in terms of impurity
 - Continue to grow the tree until you reach the leaf nodes.
- How to avoid the impact of randomization?
 - Execute this many (T) times:
 - Pick a sample Z of size N from the training data (**Bagging**)
 - Build a random tree by recursively following the steps
 - Select m random features, pick the best split, split the node
 - Output the ensemble
 - Final prediction: Average over all trees
- The first example of **ensemble of classifiers** in this course

$$p(c|\mathbf{v}) = \frac{1}{T} \sum_{t=1}^T p_t(c|\mathbf{v})$$